

In [24]: `from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from sklearn.metrics import f1_score, confusion_matrix, plot_confusion_matrix, plot_roc_curve, accuracy_score, recall_score, precision_score, balanced_accuracy_score, r
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import metrics
import seaborn as sns`

In [25]: `#load data from CSV file
data = pd.read_csv("/Users/catherinebetancourt-lee/BMEN 415/fetal_health.csv.csv")`

In [26]: `#Merging classification classes into binary
data['fetal_health'] = data['fetal_health'].replace(1.0,0)
data['fetal_health'] = data['fetal_health'].replace(2.0,0)
data['fetal_health'] = data['fetal_health'].replace(3.0,1)`

In [27]: `#Separate data and target variables
X = data.drop('fetal_health', axis = 1)
X
y = data['fetal_health']
y`

Out[27]:

0	0
1	0
2	0
3	0
4	0
..	
2121	0
2122	0
2123	0
2124	0
2125	0

Name: fetal_health, Length: 2126, dtype: int64

In [28]: `#split model into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,random_state=142)`

In [29]: `#We need to create the LASSO model
#penalty = L1 to apply the regularization, and then set saga for larger datasets.
#C = inverse of regularization strength to 0.1
lasso_model = LogisticRegression(penalty = 'l1', solver='saga', C=0.1)`

In [30]: `#fit the model into the training data
lasso_model.fit(X_train, y_train)`

/Users/catherinebetancourt-lee/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:352: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
LogisticRegression(C=0.1, penalty='l1', solver='saga')

In [31]: `#predict on testing set
y_pred = lasso_model.predict(X_test)`

In [32]: `#metrics
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
f1 = f1_score(y_test, y_pred)
print("F1 Score", f1)
balanced_acc = balanced_accuracy_score(y_test, y_pred)
print("Balanced Accuracy:", balanced_acc)
recall = recall_score(y_test, y_pred)
print("Recall Score:", recall)
precision = precision_score(y_test, y_pred)
print("Precision Score:", precision)
auc = roc_auc_score(y_test, y_pred)
print("AUC Score:", auc)
cm = confusion_matrix(y_test, y_pred)
print("Confusion matrix:")
print(cm)
print(classification_report(y_test, y_pred))`

Accuracy: 0.9647887323943662
F1 Score 0.7619047619047619
Balanced Accuracy: 0.8377420533430764
Recall Score: 0.6857142857142857
Precision Score: 0.8571428571428571
AUC Score: 0.8377420533430764
Confusion matrix:
[[387 4]
 [11 24]]

		precision	recall	f1-score	support
	0	0.97	0.99	0.98	391
	1	0.86	0.69	0.76	35
	accuracy			0.96	426
	macro avg	0.91	0.84	0.87	426
	weighted avg	0.96	0.96	0.96	426

In [33]: `sns.heatmap(cm,annot=True, annot_kws={'size':10}, fmt='d')
print(accuracy_score(y_test, y_pred))

metrics.plot_roc_curve(lasso_model, X_test, y_test)
plt.show()`

0.9647887323943662
/Users/catherinebetancourt-lee/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_roc_curve is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the class methods: :meth:`sklearn.metric.RocCurveDisplay.from_predictions` or :meth:`sklearn.metric.RocCurveDisplay.from_estimator`.
warnings.warn(msg, category=FutureWarning)

